

Faeze Ghorbanpour

faeze.ghorbanpour@tum.de • +49 176 19254896 • faezeghorbanpour.github.io
EU Blue Card Holder • German Permanent Residence • Google Scholar • LinkedIn • Github

Research Focus

Post-training and LLM alignment • parameter-efficient fine-tuning • efficient transfer learning • multilingual and multicultural NLP • model evaluation, safety, and content moderation

Experience

Doctoral Researcher, Technical University of Munich

Munich, Germany | April 2026 – Present

- Current focus: preference alignment of LLMs on subjective NLP tasks and diagnosis of multilingual dense embedders via retrieval disagreement (Under Review).
- Evaluated how LLM safety behavior changes when harmful sentences are embedded in long inputs (EMNLP 2026 Main), and how LLM judgments align with human hate-speech annotations (EMNLP 2026 Findings).

Guest Researcher, University of Copenhagen

Copenhagen, Denmark | Mar 2026 – May 2026

- With Prof. Anders Søgaard: built activation-level interventions in cross-lingual latent spaces that improve factual consistency across languages without accuracy drops (EMNLP 2026 Findings).

Doctoral Researcher, LMU Munich

Munich, Germany | Sep 2023 – Mar 2026

- Built label-specific soft-prompt and adapter methods for harmful-content detection, outperforming full fine-tuning while training under 1% of model parameters (NAACL 2025).
- Designed a cross-lingual nearest-neighbor retrieval pipeline that matches supervised performance with 50 labeled examples per language (EMNLP 2025).
- PhD started at LMU Munich; transferred to TUM in 2026 with advisor.

Guest Researcher, Max Planck Institute for Demographic Research

Rostock, Germany | Jun 2023 – Aug 2023

- Built name-based nationality inference models over 30M scholarly records to separate emigration from return migration; published at ICWSM 2025 (oral).

Research Engineer, Bimlogiq (Remote)

Sydney, Australia | Mar 2023 – Jun 2023

- Prototyped an LLM pipeline generating Autodesk Revit plugins from natural-language specs, with code validation and iterative repair.

Research Intern, Aalto University

Helsinki, Finland | Jun 2022 – Aug 2022

- Applied few-shot learning to classify communicative acts in crisis news comments; published at ASONAM 2024.

NLP Scientist, Sotoon (Hezardastan Holding)

| Jun 2019 – May 2022

- Owned Persian NLP pipelines in production (tokenization, NER, classification) serving 1000 requests/day across tens of enterprise customers.
- Shipped chatbot intent models and a content-moderation classifier at 90% F1.

NLP Scientist, Shaya Smart Solutions

| Oct 2018 – Jun 2019

- Built crawling, indexing, and entity-tracking pipelines for large-scale media monitoring.

Education

Ph.D., Computer Science, Technical University of Munich

Sep 2023 – ≈ Dec 2027

- Advisor: Prof. Alexander Fraser.
- Affiliated with the Munich Center for Machine Learning (MCML).
- Thesis: Efficient and Generalizable Transfer Learning under Resource Constraints.

M.Sc., Computer Engineering, Sharif University of Technology

Sep 2018 – Jun 2021

- Advisor: Prof. Hamid R. Rabiee.
- Thesis: similarity- and transformer-based multimodal fake news detection (published in SNAM).

B.Sc., Computer Engineering, Sharif University of Technology

Sep 2014 – Aug 2018

- Advisor: Prof. Abolfazl Motahari
- Thesis: hidden Markov models for map-matching noisy GPS traces on OpenStreetMap.

Selected Publications (* equal contribution; full list on Google Scholar)

- **Ghorbanpour, F.**, Fierro, C., Fraser, A., Søgaaard, A. Latent-Space Intervention for Cross-Lingual Factual Consistency: Consistency Improvements without Accuracy Drops. *EMNLP (Findings)*, 2026.
- Jradi, A.*, **Ghorbanpour, F.***, Fraser, A. Attribute-Based Diagnosis of LLM Alignment with Hate Speech Annotations. *EMNLP (Findings)*, 2026.
- **Ghorbanpour, F.**, Fraser, A. On the Sensitivity of LLMs to Harmful Sentences in Long Inputs. *EMNLP (Main)*, 2026.
- **Ghorbanpour, F.**, Dementieva, D., Fraser, A. Data-Efficient Hate Speech Detection via Cross-Lingual Nearest Neighbor Retrieval with Limited Labeled Data. *EMNLP (Main)*, 2025.
- **Ghorbanpour, F.**, Hangya, V., Fraser, A. Fine-Grained Transfer Learning for Harmful Content Detection through Label-Specific Soft Prompt Tuning. *NAACL (Main)*, 2025.
- **Ghorbanpour, F.**, Dementieva, D., Fraser, A. Can Prompting Unlock Hate Speech Detection across Languages? A Zero-shot and Few-shot Study. *WOAH @ ACL*, 2025.
- **Ghorbanpour, F.**, Malaguth, T., Akbaritabar, A. Differentiating Emigration from Return Migration of Scholars Using Name-Based Nationality Detection Models. *ICWSM (oral)*, 2025.
- Paakki, H., **Ghorbanpour, F.** Computational Analysis of Communicative Acts for Understanding Crisis News Comment Discourses. *ASONAM*, 2024.
- **Ghorbanpour, F.**, Ramezani, M., Fazli, M.A., Rabiee, H.R. FNR: A Similarity and Transformer-Based Approach to Detect Multi-Modal Fake News in Social Media. *Social Network Analysis and Mining*, 2023.

Preprints and Under Review

- **Ghorbanpour, F.**, Fraser, A. Preference Alignment of Language Models on Subjective NLP Tasks. 2026.
- **Ghorbanpour, F.**, Kolli, S., Diesner, J., Fraser, A. Retrieval Disagreement as a Signal: Diagnosing and Steering Multilingual Dense Embedders. 2026.
- Kolli, S., **Ghorbanpour, F.**, et al., Fraser, A., Diesner, J. What Remains Stable Across Tasks? A Multilingual, Cross-Task Evaluation of Embedding Robustness in Claim Matching. 2026.

Technical Skills

ML & NLP: PyTorch, Hugging Face Transformers, PEFT/LoRA, Adapters, TRL, Unsloth, vLLM, Weights & Biases

Programming Languages: Python, Java, C/C++, R, Bash

Training and serving: multi-GPU training (DeepSpeed, torchrun), SLURM clusters, quantized inference, FastAPI, Docker

Retrieval and data: FAISS, ElasticSearch, PySpark, SQL (PostgreSQL, MySQL), MongoDB

Awards and Service

- Reviewer, ACL Rolling Review, 2026 – present.
- MCML X-Change research visit grant, University of Edinburgh, 2026.
- TUM Graduate School research visit grant, University of Copenhagen, 2026.
- Ranked 13th of 824 candidates, national AI Ph.D. entrance exam, 2020.

Teaching and Supervision

- Supervised 9 BSc/MSc theses and 4 Praktikum teams on fairness, LLM agents, and evaluation, TU Munich, 2024 – present.
- Teaching assistant, LLM and Information Extraction seminars, LMU Munich, 2023 – 2024.
- Tutor, Advanced Machine Learning, Information Retrieval, and Databases, Sharif University of Technology, 2018 – 2020.

Languages

English (fluent), German (A2), Turkish (basic), Persian and Azerbaijani (native)